

example_2

July 19, 2019

1 Example 2 - Combining Machine Learning and Operations Research methods to advance the Project Management Practice

```
In [1]: import json
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB

import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
nltk.download('wordnet')
from nltk.stem import WordNetLemmatizer
wordnet_lemmatizer = WordNetLemmatizer()

from nltk.stem import LancasterStemmer
stemmer = LancasterStemmer()

from lime.lime_text import LimeTextExplainer

from ortools.graph import pywrapgraph

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

%matplotlib inline
```

```
[nltk_data] Downloading package stopwords to /home/nkanak/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /home/nkanak/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
In [2]: # Data has been retrieved from the public accessible Jira instance of the open source project
# Check this URL: https://issues.apache.org/jira .
```

```

with open('data/hadoop_issues.json') as f:
    issues = json.load(f)
    issues = issues['issues']

In [3]: extra_stopwords = ["a", "about", "above", "after", "again", "against", "ain", "all", "an", "and", "are", "as", "at", "be", "because", "before", "but", "by", "can", "could", "do", "each", "every", "for", "from", "has", "have", "he", "her", "his", "hundred", "if", "in", "into", "is", "it", "its", "me", "more", "most", "my", "of", "off", "on", "one", "only", "or", "other", "out", "over", "so", "some", "than", "that", "the", "there", "these", "they", "this", "those", "to", "too", "two", "us", "was", "were", "we", "which", "while", "who", "with", "without", "would", "you"]

In [4]: # Keep only the issues with assignee.
    issues = [issue for issue in issues if issue['fields'].get('assignee') is not None]

In [5]: len(issues)

Out[5]: 677

In [6]: # Keep only the issues of the 4 most important employees, i.e. employees with the highest number of issues.
    issues = [issue for issue in issues if issue['fields']['assignee']['key'] in ['stevel@apache.org', 'gabor.bota', 'danielzhou', 'ajisakaa']]

In [7]: fake_names = {
    'stevel@apache.org': 'john',
    'gabor.bota': 'jane',
    'danielzhou': 'johnny',
    'ajisakaa': 'richard'
}
for issue in issues:
    issue['fields']['assignee']['key'] = fake_names[issue['fields']['assignee']['key']]

In [8]: len(issues)

Out[8]: 174

In [9]: unique_assignees_to_number_mapping = {assignee: key for key, assignee in enumerate(sorted(unique_person_names))}
    unique_person_names = sorted([key for key in unique_assignees_to_number_mapping], key=lambda key: unique_assignees_to_number_mapping[key])
    print(unique_assignees_to_number_mapping)

{'john': 0, 'jane': 1, 'johnny': 2, 'richard': 3}

In [10]: columns = {
    'class': [issue['fields']['assignee']['key'] for issue in issues],
    'text': [(issue['fields']['description'] if issue['fields']['description'] is not None else '') for issue in issues]
}

In [11]: # Compose and clean up text.
    # The text of each issue is composed of two attributes, namely description and summary.
    # Removal of english stop words is performed as well as lemmatization and stemming.
    # Lemmatization (a Text Normalization technique) is the process of grouping together words that have the same root.
    # Stemming is the process of reducing inflected (or sometimes derived) words to their root form.
    # Lemmatization, unlike Stemming, reduces the inflected words properly ensuring that the root form is the same.
    # Stemming is different to Lemmatization in the approach it uses to produce root forms.
    # Also lemmatization and stemming techniques decrease the number of features of each document.
    for i in range(len(columns['text'])):
        columns['text'][i] = ' '.join([stemmer.stem(wordnet_lemmatizer.lemmatize(word.lower())) for word in columns['text'][i].split()])

```

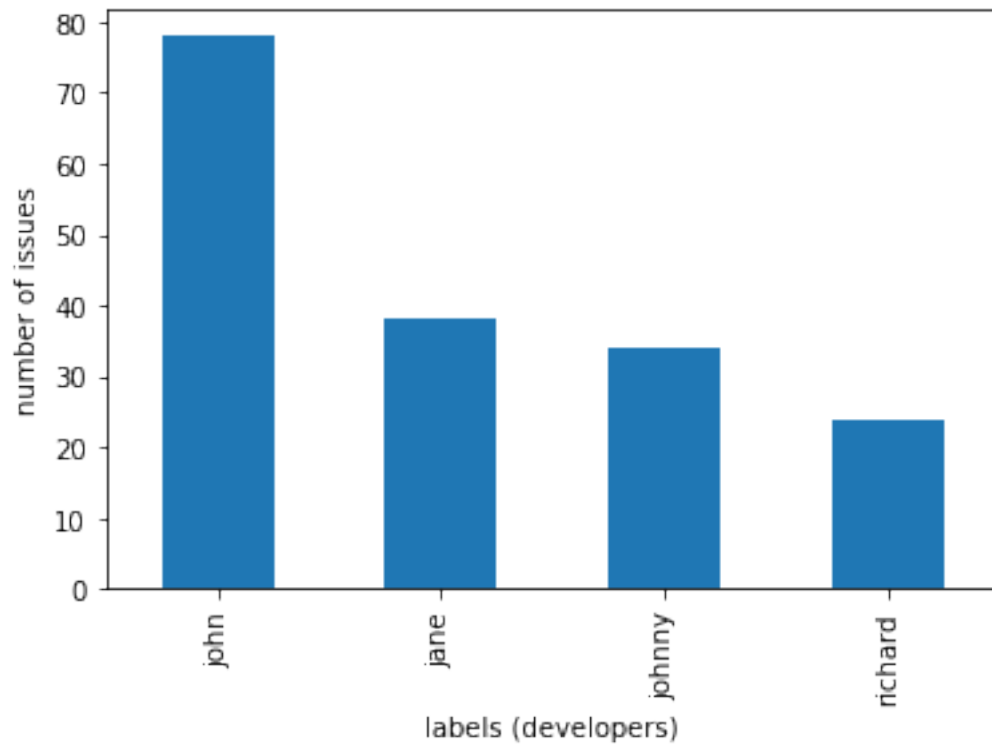
```
In [12]: issues_df = pd.DataFrame.from_dict(columns)
         #issues_df
```

```
In [13]: issues_df['class'].value_counts()
```

```
Out[13]: john      78
         jane      38
         johnny    34
         richard   24
         Name: class, dtype: int64
```

```
In [14]: ax = issues_df['class'].value_counts().plot(kind='bar')
         ax.set_xlabel('labels (developers)')
         ax.set_ylabel('number of issues')
```

```
Out[14]: Text(0, 0.5, 'number of issues')
```



1.1 Example

```
In [15]: #issues_df
```

```
In [16]: issues_df['label'] = issues_df['class'].map(unique_assignees_to_number_mapping)
```

```
In [17]: X = issues_df['text']
y = issues_df['label']
# Tokenization, tag/feature generation and removal of frequently occurred tags/words.
cv = CountVectorizer(stop_words='english', max_df=0.6)

X_transformed = cv.fit_transform(X)
X_train, X_test, y_train, y_test, indices_train, indices_test = train_test_split(X_train, X_test, y_train, y_test, indices_train, indices_test)
print('Train size %s' % X_train.shape[0])
print('Test size %s' % X_test.shape[0])
print('Number of features %s' % X_test.shape[1])
```

Train size 116

Test size 58

Number of features 2403

```
In [18]: # Naive Bayes classifier for multinomial models.
# The multinomial Naive Bayes classifier is suitable for classification with discrete
# The multinomial distribution normally requires integer feature counts. However, in
naive_clf = MultinomialNB()
naive_clf.fit(X_train, y_train)
naive_clf.score(X_test, y_test)
```

Out[18]: 0.8275862068965517

```
In [19]: y_predicted = [unique_person_names[i] for i in naive_clf.predict(X_transformed[indices_test])]
y_true = issues_df.iloc[indices_train]['class'].to_list()
class_names = unique_person_names
```

```
def plot_confusion_matrix(y_true, y_pred, classes,
                           normalize=False,
                           title=None,
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if not title:
        if normalize:
            title = 'Normalized confusion matrix'
        else:
            title = 'Confusion matrix, without normalization'

    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    # Only use the labels that appear in the data
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
```

```

fig, ax = plt.subplots()
im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
ax.figure.colorbar(im, ax=ax)
# We want to show all ticks...
ax.set(xticks=np.arange(cm.shape[1]),
       yticks=np.arange(cm.shape[0]),
       # ... and label them with the respective list entries
       xticklabels=classes, yticklabels=classes,
       title=title,
       ylabel='True label',
       xlabel='Predicted label')

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
         rotation_mode="anchor")

# Loop over data dimensions and create text annotations.
fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, format(cm[i, j], fmt),
                ha="center", va="center",
                color="white" if cm[i, j] > thresh else "black")
fig.tight_layout()
return ax

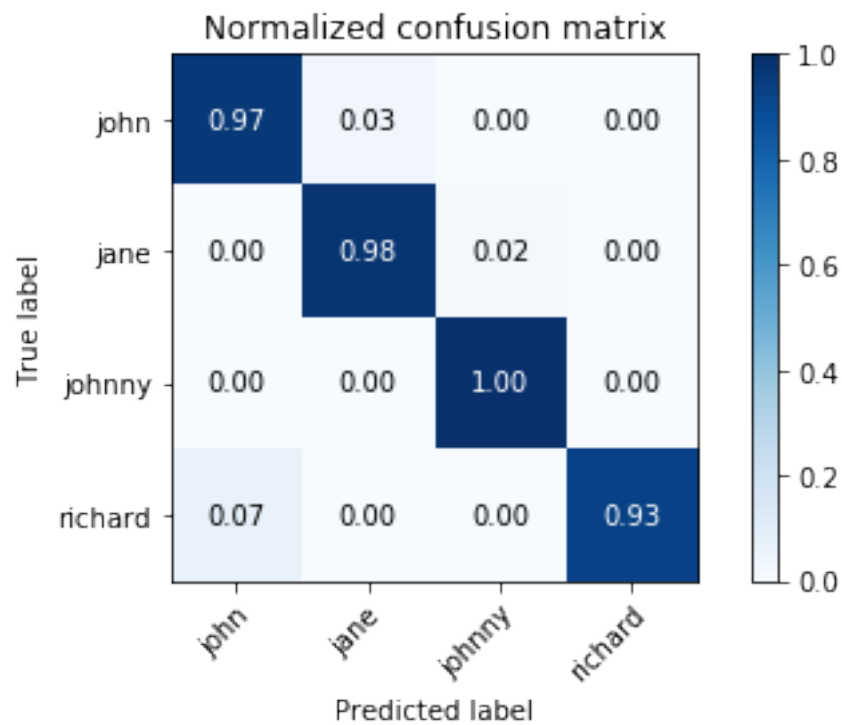
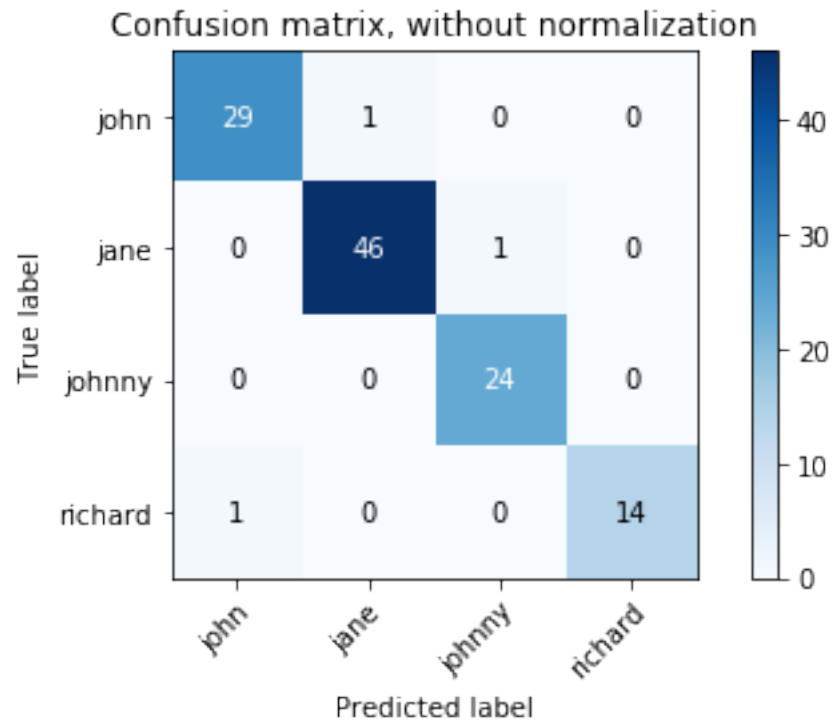
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plot_confusion_matrix(y_true, y_predicted, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plot_confusion_matrix(y_true, y_predicted, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()

```



```

In [20]: # Text Explainer for explaining the selected examples.
# Reference: https://arxiv.org/abs/1602.04938
# The Explanations help us to check the reliability and validity of the trained machine
# The Explanations confirm that the model chooses the right label/class for the right
explainer = LimeTextExplainer(class_names=unique_person_names)

In [21]: def explain_classification(text, classifier):
    return explainer.explain_instance(text, lambda x: classifier.predict_proba(cv.transform(x)))

In [22]: unique_assignees_to_number_mapping

Out[22]: {'john': 0, 'jane': 1, 'johnny': 2, 'richard': 3}

In [23]: def analyze_selected_examples(index):
    print(index)
    print('Real selected label: %s' % issues_df.iloc[index]['class'])
    print('Probabilities of each label: %s' % naive_clf.predict_proba(cv.transform([issues_df.iloc[index]['text']])))
    print('Summary: %s' % issues[index]['fields']['summary'])
    print('Description: %s' % issues[index]['fields']['description'] if issues[index]['fields']['description'] else '')
    exp = explain_classification(issues_df.iloc[index]['text'], naive_clf)
    exp.show_in_notebook()
    exp.save_to_file('example_explanations/%s.html' % (index))
    return exp

In [24]: analyze_selected_examples(138)

138
Real selected label: john
Probabilities of each label: [[9.63e-01 7.62e-11 4.10e-08 3.70e-02]]
Summary: branch-2 site not building after ADL troubleshooting doc added
Description: Toc error on the ADL troubleshooting doc from HADOOP-15090
{code}
[ERROR] Failed to execute goal org.apache.maven.plugins:maven-site-plugin:3.5:site (default-cli) on project maven-site-plugin:3.5:site (default-cli): The goal has failed
{code}

<IPython.core.display.HTML object>

Out[24]: <lime.explanation.Explanation at 0x7f77c65a76a0>

In [25]: analyze_selected_examples(65)

65
Real selected label: john
Probabilities of each label: [[9.49e-01 1.43e-02 3.67e-02 8.99e-08]]
Summary: S3 listing inconsistency can raise NPE in globber
Description: FileSystem Globber does a listStatus(path) and then, if only one element is returned, it returns the first element. This is not correct as it should return the first element that is not a directory.

```

On S3, if the path has had entries deleted, the listing may include files which are no longer t

While its wrong to glob against S3 when its being inconsistent, we should at least fail gracef

Proposed

```
# log all IOEs raised in Globber.getFileStatus @ debug
# catch FNFES and downgrade to warn
# continue
```

The alternative would be fail fast on FNFE, but that's more traumatic

<IPython.core.display.HTML object>

Out[25]: <lime.explanation.Explanation at 0x7f77c1adea58>

In [26]: analyze_selected_examples(85)

85

Real selected label: johnny

Probabilities of each label: [[8.42e-02 2.02e-04 9.16e-01 3.79e-07]]

Summary: ABFS: Code changes for bug fix and new tests

Description: - add bug fixes.

- remove unnecessary dependencies.

- add new tests for code changes.

<IPython.core.display.HTML object>

Out[26]: <lime.explanation.Explanation at 0x7f77c6db4d68>

In [27]: analyze_selected_examples(109)

109

Real selected label: john

Probabilities of each label: [[0.53 0.43 0.01 0.03]]

Summary: Release Hadoop 2.7.7

Description: Time to get a new Hadoop 2.7.x out the door.

<IPython.core.display.HTML object>

Out[27]: <lime.explanation.Explanation at 0x7f77c1a876d8>


```

In [28]: # OR part of the example.
        # Target?: To maximize the chance of success of the software company.
        # How?: By assigning the employees in such a way that the total relevance is maximized.
        # Relevance = skills required by an issue vs skills possessed by an employee.
        # In our example the 'relevance' is equivalent to the probability (calculated by our
        # OR algorithm?: Linear Assignment Problem (LAP) (https://developers.google.com/optim

In [29]: selected_example_indeces = [138, 65, 85, 109]

In [30]: # Row == employee.
        # Column == issue.
        # Transpose a list code: list(map(list, zip(*l)))
relevance_of_each_employee_per_issue = list(map(list, zip(*[
    list(naive_clf.predict_proba(cv.transform([issues_df.iloc[i]['text']])))[0]) for i in
    selected_example_indeces])))
relevance_of_each_employee_per_issue_percentage = [[int(round(c*100)) for c in row] for row in
    relevance_of_each_employee_per_issue]
print('Relevance of each employee per issue: %s' % relevance_of_each_employee_per_issue_percentage)

Relevance of each employee per issue: [[96, 95, 8, 53], [0, 1, 0, 43], [0, 4, 92, 1], [4, 0, 0, 100]]

In [31]: def assign_employees_to_issues():
        cost = create_data_array()
        rows = len(cost)
        cols = len(cost[0])

        assignment = pywrapgraph.LinearSumAssignment()
        for worker in range(rows):
            for task in range(cols):
                if cost[worker][task]:
                    assignment.AddArcWithCost(worker, task, cost[worker][task])
        solve_status = assignment.Solve()
        if solve_status == assignment.OPTIMAL:
            total_relevance = 0
            for i in range(0, assignment.NumNodes()):
                relevance = relevance_of_each_employee_per_issue_percentage[i][assignment.RightMate(i)]
                total_relevance += relevance
                print('Employee %s (index:%s) is assigned to issue %s. Relevance = %d' % (
                    unique_person_names[i],
                    i,
                    selected_example_indeces[assignment.RightMate(i)],
                    relevance))
            print()
            print('Total relevance = ', total_relevance)
        elif solve_status == assignment.INFEASIBLE:
            print('No assignment is possible.')
        elif solve_status == assignment.POSSIBLE_OVERFLOW:
            print('Some input costs are too large and may cause an integer overflow.')

```

```

def create_data_array():
    cost = relevance_of_each_employee_per_issue_percentage
    inverse_cost = [[100 - c for c in row] for row in cost]
    #print(cost)
    #print(inverse_cost)
    return inverse_cost

```

In [32]: # The Outcome of the example.

```

# Evaluation http://www.hungarianalgorithm.com/solve.php?c=96-95-8-53--0-4-92-1--4-0-
assign_employees_to_issues()

```

Employee john (index:0) is assigned to issue 65. Relevance = 95
 Employee jane (index:1) is assigned to issue 109. Relevance = 43
 Employee johnny (index:2) is assigned to issue 85. Relevance = 92
 Employee richard (index:3) is assigned to issue 138. Relevance = 4

Total relevance = 234